

# Intermediate Machine Learning: Assignment 4

## Deadline

Assignment 4 is due Monday, November 18 by 11:59pm. Late work will not be accepted as per the course policies (see the Syllabus and Course policies on Canvas).

Directly sharing answers is not okay, but discussing problems with the course staff or with other students is encouraged.

You should start early so that you have time to get help if you're stuck. The drop-in office hours schedule can be found on Canvas. You can also post questions or start discussions on Ed Discussion. The assignment may look long at first glance, but the problems are broken up into steps that should help you to make steady progress.

## Submission

Submit your assignment as a pdf file on Gradescope, and as a notebook (.ipynb) on Canvas. You can access Gradescope through Canvas on the left-side of the class home page. The problems in each homework assignment are numbered. Note: When submitting on Gradescope, please select the correct pages of your pdf that correspond to each problem. This will allow graders to more easily find your complete solution to each problem.

To produce the .pdf, please do the following in order to preserve the cell structure of the notebook:

Go to "File" at the top-left of your Jupyter Notebook Under "Download as", select "HTML (.html)" After the .html has downloaded, open it and then select "File" and "Print" (note you will not actually be printing) From the print window, select the option to save as a .pdf

## Topics

- Graph kernels
- Reinforcement learning

This assignment will also help to solidify your Python skills.

## Problem 1: Graph kernels (20 points)

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as img
```

```
import sklearn
import random
from numpy.linalg import inv
%matplotlib inline
```

In [2]: *# Helper functions for third part of exercise*

```
def rgb2gray(rgb):
    """Function to turn RGB images in greyscale images."""
    return np.dot(rgb[..., :3], [0.2989, 0.5870, 0.1140])

def grid_adj(rows, cols):
    """Function that creates the adjacency matrix of
    a grid graph with predefined amount of rows and columns."""
    M = np.zeros([rows*cols, rows*cols])
    for r in np.arange(rows):
        for c in np.arange(cols):
            i = r*cols + c
            if c > 0:
                M[i-1, i] = M[i, i-1] = 1
            if r > 0:
                M[i-cols, i] = M[i, i-cols] = 1
    return M
```

The graph Laplacian for a weighted graph on  $n$  nodes is defined as

$$L = D - W$$

where  $W$  is an  $n \times n$  symmetric matrix of positive edge weights, with  $W_{ij} = 0$  if  $(i, j)$  is not an edge in the graph, and  $D$  is the diagonal matrix with  $D_{ii} = \sum_{j=1}^n W_{ij}$ . This generalizes the definition of the Laplacian used in class, where all of the edge weights are one.

1. Show that  $L$  is a Mercer kernel, by showing that  $L$  is symmetric and positive-semidefinite.
2. In graph neural networks we define polynomial filters of the form

$$P = a_0 I + a_1 L + a_2 L^2 + \dots a_d L^d$$

where  $L$  is the Laplacian and  $a_0, \dots, a_d$  are parameters, corresponding to the filter parameters in standard convolutional neural networks.

If each  $a_i \geq 0$  is non-negative, show that  $P$  is also a Mercer kernel.

3. This polynomial filter has many applications. A handful of these applications are based on the fact that, given a graph with a signal  $x$ , the value of  $x^T L x$  will be low in case the signal is smooth (i.e. smooth transitions of  $x$  between neighboring nodes). A large  $x^T L x$  means that we have a rough graph signal (i.e. a lot of jumps in  $x$  between neighboring nodes).

An interesting application that uses this property is the so-called image inpainting process, where an image is seen as grid graph. Image inpainting tries to restore a

corrupted image by smoothing out the neighboring pixel values. In this problem we corrupt an image by turning off (i.e. making the pixel value equal to zero) a certain portion of the pixels. Your goal will be to restore the corrupted image and hence recreate the original image.

First, let's corrupt an image by turning off a portion of the pixels. For this exercise, we choose to turn off 30% of the pixels. The result is shown below. Try to understand the code, as some variables might be interesting for your work.

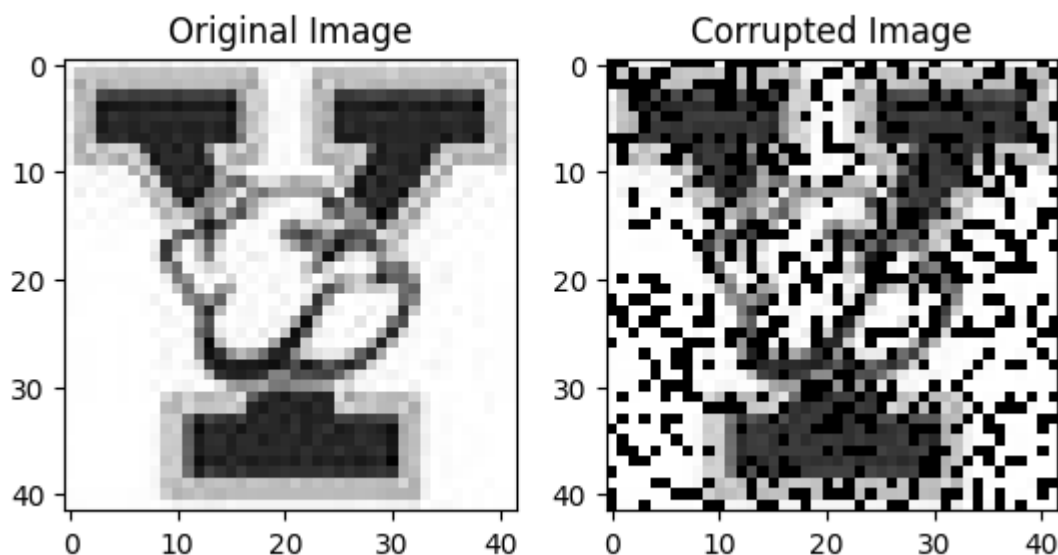
The image "Yale\_Bulldogs.jpg" can be found in Canvas under assn4 folder, and also in the GitHub repo <https://github.com/YData123/sds365-fa24/tree/main/assignments/assn4>

```
In [4]: # Normalize the pixels of the original image
image = img.imread("Yale_Bulldogs.jpg") / 255
# Turn picture into greyscale
gray_image = rgb2gray(image)
height_img = gray_image.shape[0]
width_img = gray_image.shape[1]

# Turn off (value 0) certain pixels
fraction_off = int(0.30*height_img*width_img)
mask = np.ones(height_img*width_img, dtype=int)
# Set the first fraction of pixels off
mask[:fraction_off] = 0
# Shuffle to create randomness
np.random.shuffle(mask)
# Multiply the original image by the reshapes mask
mask = np.reshape(mask, (height_img, width_img))
corrupted_image = np.multiply(mask, gray_image)

fig, (ax1, ax2) = plt.subplots(1, 2)
ax1.imshow(gray_image, cmap=plt.get_cmap('gray'))
ax1.set_title("Original Image")
ax2.imshow(corrupted_image, cmap=plt.get_cmap('gray'))
ax2.set_title("Corrupted Image")

plt.show()
```



Inpainting missing pixel values can be formulated as the following optimization problem:

$$\min_{\mathbf{x} \in \mathbb{R}^n} \{ \|\mathbf{y} - \mathbf{M}\mathbf{x}\|_2^2 + \alpha \mathbf{x}^T \mathbf{P} \mathbf{x} \}$$

where  $\mathbf{y} \in \mathbb{R}^n$  ( $n$  being the total amount of pixels) is the corrupted graph signal (with missing pixel values being 0) and  $\alpha$  is a regularization (smoothing) parameter that controls for smoothness of the graph.  $\mathbf{P}$  is the polynomial filter based on the laplacian  $\mathbf{L}$ . Finally,  $\mathbf{M} \in \mathbb{R}^{n \times n}$  is a diagonal matrix that satisfies:

$$\mathbf{M}(i, i) = \begin{cases} 1, & \text{if } \mathbf{y}(i) \text{ is observed} \\ 0, & \text{if } \mathbf{y}(i) \text{ is corrupted} \end{cases}$$

The optimization problem tries to find an  $\mathbf{x}$  that matches the observed values in  $\mathbf{y}$ , and at the same time tries to be smooth on the graph. Start with deriving a closed form solution of this optimization problem:

## your markdown here

To solve this optimization problem, we can use the method of Lagrange multipliers to derive a closed-form solution. The objective function is given by:

$$f(\mathbf{x}) = \|\mathbf{y} - \mathbf{M}\mathbf{x}\|_2^2 + \alpha \mathbf{x}^T \mathbf{P} \mathbf{x}$$

where  $(\textbf{M})$  is a diagonal matrix that differentiates between missing and observed pixels. We start by expanding the objective function:

$$f(\mathbf{x}) = (\mathbf{y} - \mathbf{M}\mathbf{x})^T (\mathbf{y} - \mathbf{M}\mathbf{x}) + \alpha \mathbf{x}^T \mathbf{P} \mathbf{x}$$

Expanding and simplifying this, we get:

$$f(\mathbf{x}) = \mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{M}\mathbf{x} + \mathbf{x}^T \mathbf{M}^T \mathbf{M}\mathbf{x} + \alpha \mathbf{x}^T \mathbf{P} \mathbf{x}$$

To find the optimal  $(\textbf{x})$ , we take the derivative with respect to  $(\textbf{x})$  and set it to zero:

$$-2\mathbf{M}^T \mathbf{y} + 2\mathbf{M}^T \mathbf{M}\mathbf{x} + 2\alpha \mathbf{P} \mathbf{x} = 0$$

This simplifies to:

$$(\mathbf{M}^T \mathbf{M} + \alpha \mathbf{P}) \mathbf{x} = \mathbf{M}^T \mathbf{y}$$

Since  $(\textbf{M})$  is a diagonal matrix (with entries of either 0 or 1),  $(\textbf{M})^T = \textbf{M}$ . Therefore, we can rewrite the equation as:

$$(\mathbf{M} + \alpha \mathbf{P}) \mathbf{x} = \mathbf{M} \mathbf{y}$$

Thus, the closed-form solution for  $(\textbf{x})$  is:

$$\mathbf{x} = (\mathbf{M} + \alpha \mathbf{P})^{-1} \mathbf{M} \mathbf{y}$$

This expression provides the restored image by taking into account both the observed values in  $\mathbf{y}$  and smoothing based on the polynomial filter  $\mathbf{P}$ . The regularization parameter  $\alpha$  controls the trade-off between fitting the observed data and promoting smoothness in the reconstructed image.

Next, let's restore our image. To keep things simple, let's say we already trained the polynomial filter  $\mathbf{P}$  of degree 2 and we found the following weights:

$$\mathbf{P} = \mathbf{L} + 0.05 \mathbf{L}^2$$

Fill in the following lines of code and show your reconstructed images next to the corrupted image. Assume that the weights on the graph edges are equal to 1.

```
In [5]: # Define the corrupted graph signal
y = corrupted_image.flatten()

# Define the diagonal matrix M based on observed pixels
M = np.diag(mask.flatten())

# Create the adjacency matrix A of the graph
A = grid_adj(height_img, width_img)

# Define the degree matrix D
D = np.diag(A.sum(axis=1))

# Define the graph Laplacian L
L = D - A


# Polynomial filter P (for example, using a simple form like P = L)
# You can adjust coefficients for higher-order terms if desired
a_0, a_1 = 1, 1
P = a_0 * np.eye(L.shape[0]) + a_1 * L

In [6]: # closed form solution you derived above
# Solve for x using the closed-form solution with different alpha values
alpha = 0.1 # Adjust alpha for different smoothing effects
x = inv(M + alpha * P).dot(M).dot(y)

In [7]: # Try to experiment with different alpha values
alpha = 0.1
reconstructed_image = np.reshape(x, (height_img, width_img))

fig, (ax1, ax2) = plt.subplots(1, 2)
ax1.imshow(corrupted_image, cmap=plt.get_cmap('gray'))
ax1.set_title("Corrupted Image")
ax2.imshow(reconstructed_image, cmap=plt.get_cmap('gray'))
ax2.set_title("Reconstructed Image")

plt.show()
```



4. Discuss the influence of the smoothing parameter  $\alpha$  in the optimization problem above. What happens for very large and very low values of  $\alpha$ ? Finally, discuss the degree of our polynomial function  $\mathbf{P}$ . What happens if we would choose a large degree?

### Influence of $\alpha$ :

#### 1. Very Large $\alpha$ :

When  $\alpha$  is large, the regularization term  $\alpha \mathbf{x}^T \mathbf{P} \mathbf{x}$  becomes dominant in the objective function. This means the solution will prioritize smoothness over fidelity to the original image data. As a result, the reconstructed image will appear overly smoothed, potentially losing important details and high-frequency information, which could make edges and textures in the image less distinguishable. In extreme cases, the reconstructed image might become nearly uniform, lacking any contrast or texture.

#### 2. Very Low $\alpha$ :

When  $\alpha$  is close to zero, the data-fitting term  $\|\mathbf{y} - \mathbf{M}\mathbf{x}\|_2^2$  becomes more influential. This leads to a solution that closely matches the observed values in  $\mathbf{y}$  but might not enforce smooth transitions between pixels. Consequently, the reconstructed image will better retain original details but may exhibit abrupt changes or noise, especially in regions where data is missing. This could result in a "noisy" appearance in the reconstructed image, with high-frequency artifacts at the edges of missing data regions.

### Influence of the Polynomial Degree of $\mathbf{P}$ :

The degree of the polynomial function  $\mathbf{P}$  influences the filter's ability to capture different frequencies in the graph signal, which affects how smooth the reconstructed image will be.

#### 1. Low Degree of $\mathbf{P}$ :

With a low degree (e.g.,  $d = 1$  or  $d = 2$ ),  $\mathbf{P}$  mainly captures low-frequency components of the graph signal. This degree will only enforce basic smoothness, resulting in a moderate level of smoothing, which might be insufficient for effectively diffusing information across larger gaps in missing data.

## 2. High Degree of $\mathbf{P}$ :

As the degree of  $\mathbf{P}$  increases, the filter becomes more capable of capturing higher frequencies in the graph signal. This can help propagate information across more distant nodes in the graph, resulting in a stronger smoothing effect. However, if the degree is too high, it could lead to over-smoothing and loss of detail, similar to the effect of a very large *alpha*. Additionally, a high polynomial degree increases the computational cost and can lead to numerical instability, especially in large graphs.

# Problem 2: Positive reinforcement (10 points)

As discussed in class, reinforcement learning using policy gradient methods is based on maximizing the expected total reward

$$J(\theta) = \mathbb{E}_{\theta}[R(\tau)],$$

where the expectation is over the probability distribution over sequences  $\tau$  through a choice of actions using the policy. This can be rewritten as

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\theta} [R(\tau) \nabla_{\theta} \log p(\tau | \theta)].$$

Approximating this gradient involves computing  $\nabla_{\theta} \log \pi_{\theta}(a | s)$  where  $\pi_{\theta}$  is the policy.

## 2.1 Continuous action space with Gaussian policy

Suppose that the action space is continuous and  $\pi_{\theta}(a | s)$  is a normal density with mean  $\mu_{\theta}(s)$  and variance  $\sigma_{\theta}^2(s)$ , two outputs of a neural network with input  $s$  and parameters  $\theta$ .

Suppose the outputs of the neural network are given by

$$\begin{aligned} \mu_{\theta}(s) &= \beta_1^T h(s) \\ \sigma_{\theta}^2(s) &= \exp(\beta_2^T h(s)) \end{aligned}$$

where  $h(s)$  is the vector of neurons in the last layer, immediately before the outputs. Derive explicit expressions for  $\nabla_{\beta_1} \log \pi_{\theta}(a | s)$  and  $\nabla_{\beta_2} \log \pi_{\theta}(a | s)$ .

Explain how these gradients and other gradient terms in  $\nabla_{\theta} \log \pi_{\theta}(a | s)$  are used to estimate the policy.

## 2.2 Discrete action space with Softmax policy


Suppose the action space is discrete with  $K$  possible actions, and the policy  $\pi_\theta(a \mid s)$  is defined using a softmax function over preferences  $u_\theta(s, a)$ :

$$\pi_\theta(a \mid s) = \frac{\exp(u_\theta(s, a))}{\sum_a \exp(u_\theta(s, a))},$$

where  $u_\theta(s, a) = \beta^T h(s, a)$ , and  $h(s, a)$  is a feature vector for state-action pair  $(s, a)$ . Derive the expression for  $\nabla_\beta \log \pi_\theta(a \mid s)$ .

## Problem 3: Deep Q-Learning for Flappy Bird (25 points)

Deep Q-learning was proposed (and patented) by DeepMind and made a big splash when the same deep neural network architecture was shown to be able to surpass human performance on many different Atari games, playing directly from the pixels. In this problem, we will walk you through the implementation of deep Q-learning to learn to play the Flappy Bird game.



The implementation is based these references:

- [DeepLearningFlappyBird](#)
- [Deep Q-Learning for Atari Breakout](#)

We use the `pygame` package to visualize the interaction between the algorithm and the game environment. However, `pygame` is not well supported by Google Colab; we recommend you to run the code for this problem locally. A window will be popped up that displays the game as it progress in real-time (as for the Cartpole demo from class).

This problem is structured as follows:

- Load necessary packages
- Test the visualization of the game, to make sure everything's working
- Process the images to reduce the dimension



- Setup the game history buffer
- Implement the core Q-learning function
- Run the learning algorithm
- Interpret the results

## Introduction

The Flappy Bird game is requires a few Python packages. Please install these *as soon as possible*, and notify us of any issues you experience so that we can help. The Python files can also be found on Canvas and in our GitHub repo at <https://github.com/YData123/sds365-fa24/tree/main/assignments/assn4>.

```
In [31]: # %pip install pygame
# %pip install opencv-python
import numpy as np
import cv2
import wrapped_flappy_bird as flappy_bird
from collections import deque
import random
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, initializers
```

## The Flappy Bird environment

Interaction with the game environment is carried out through calls of the form

```
(image, reward, terminal) = game.frame_step(action)
```

where the meaning of these variables is as follows:

- `action` :  $\begin{pmatrix} 1 \\ 0 \end{pmatrix}$  for doing nothing,  $\begin{pmatrix} 0 \\ 1 \end{pmatrix}$  for "flapping the bird's wings"
- `image` : the image for the next step of the game, of size (288, 512, 3) with three RGB channels
- `reward` : the reward received for taking the action; -1 if an obstacle is hit, 0.1 otherwise.
- `terminal` : True if an obstacle is hit, otherwise False

Now let's take a look at the game interface. First, initiate the game:

```
In [32]: num_actions = 2

# initiate a game
game = flappy_bird.GameState()

# get the first state by doing nothing
do_nothing = np.zeros(num_actions)
do_nothing[0] = 1
image, reward, terminal = game.frame_step(do_nothing)

print('shape of image:', image.shape)
print('reward: ', reward)
```

```
print('terminal: ', terminal)
```

shape of image: (288, 512, 3)

reward: 0.1

terminal: False

After running the above cells, a window should pop up, and you can watch the game being played in that window.


Let's take some random actions and see what happens:

```
In [33]: for i in range(587):  
  
    # choose a random action  
    action = np.random.choice(num_actions)  
  
    # create the corresponding one-hot vector  
    action_vec = np.zeros(num_actions)  
    action_vec[action] = 1  
  
    # take the action and observe the reward and the next state  
    image, reward, terminal = game.frame_step(action_vec)
```

Are you able to see Flappy moving across the window and crashing into things? Great! If you're having any issues, post to EdD and we'll do our best to help you out.

Here is how we can visualize a frame of the game as an image within a cell.

```
In [34]: # show the image  
import matplotlib.pyplot as plt  
plt.imshow(image.transpose([1, 0, 2]))  
plt.show()  
plt.close()
```



## Preprocessing the images


Alright, next we need to preprocess the images by converting them to grayscale and resizing them to  $80 \times 80$  pixels. This will help to reduce the computation, and aid learning. Besides, Flappy is "color blind." (Fun fact: The instructor of this course is also [color vision deficient](#).)

```
In [35]: def resize_gray(frame):
        frame = cv2.cvtColor(cv2.resize(frame, (80, 80)), cv2.COLOR_BGR2GRAY)
        ret, frame = cv2.threshold(frame, 1, 255, cv2.THRESH_BINARY)
        return np.reshape(frame, (80, 80, 1))

        image_transformed = resize_gray(image)
        print('Shape of the transformed image:', image.shape)

        # show the transformed image
        _ = plt.imshow(image_transformed.transpose((1, 0, 2)), cmap='gray')
```

Shape of the transformed image: (288, 512, 3)



This shows the preprocessed image for a single frame of the game. In our implementation of Deep Q-Learning, we encode the state by stacking four consecutive frames, resulting in a tensor of shape (80,80,4).

Then, given the `current_state`, and a raw image `image_raw` of size  $288 \times 512 \times 3$ , we convert the raw image to a  $80 \times 80 \times 1$  grayscale image using the code in the previous cell. Then, we remove the first frame of `current_state` and add the new frame, giving again a stack of images of size (80, 80, 4).

```
In [36]: def preprocess(image_raw, current_state=None):
        # resize and convert to grayscale
```

```
image = resize_gray(image_raw)
# stack the frames
if current_state is None:
    state = np.concatenate((image, image, image, image), axis=2)
else:
    state = np.concatenate((image, current_state[:, :, :3]), axis=2)
return state
```

### 3.1 Explain the game state

Why is the state chosen to be a stack of four consecutive frames rather than a single frame? Give an intuitive explanation:

## answer:

**let agent have memory: to determine whether objects are moving closer or farther away, or to estimate the agent's velocity and acceleration relative to other objects.**

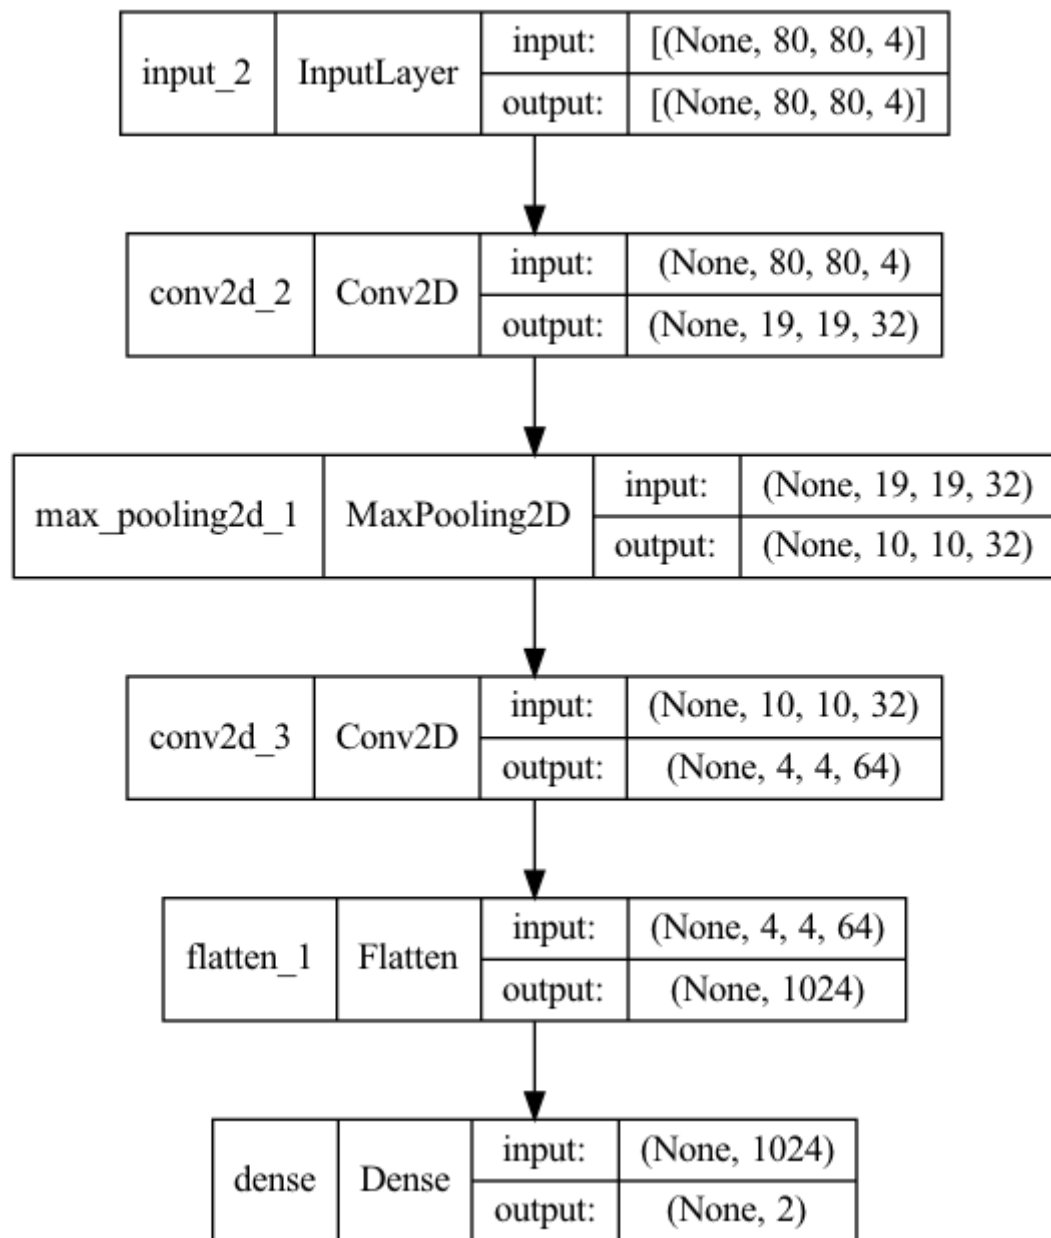
The state is chosen as a stack of four consecutive frames rather than a single frame to capture temporal information about the game environment. This is particularly useful in environments where movement and object dynamics are key, where the agent needs to understand the direction and speed of obstacles (like pipes) and itself.

If we use only a single frame, we lose the temporal context, making it difficult for the agent to determine whether objects are moving closer or farther away, or to estimate the agent's velocity and acceleration relative to other objects. By stacking four frames, the agent gains insight into these temporal aspects, allowing it to learn better strategies for navigation and obstacle avoidance based on the movement history contained in those frames. This temporal awareness is crucial for making accurate decisions in real-time.

### Constructing the neural network

Now we are ready to construct the neural network for approximating the Q function. Recall that, given input  $s$  which is of size  $80 \times 80 \times 4$  due to the previous preprocessing, the output of the network should be of size 2, corresponding to the values of  $Q(s, a_1)$  and  $Q(s, a_2)$  respectively.

Here is the summary of the model we'd like to build:



## 3.2 Initialize the network

Complete the code in the next cell so that your model architecture matches that in the above picture. Here we specify the initialization of the weights by using `keras.initializers`. Note that we haven't talked about the `strides` argument for CNNs; you can read about stride here: <https://machinelearningmastery.com/padding-and-stride-for-convolutional-neural-networks/>. It's not important to understand this in detail, you just need to choose the number and sizes of the filters to get the shapes to match the specification.

```
In [42]: from tensorflow.keras import layers, initializers
from tensorflow import keras

def create_q_model():
    # Input layer
    state = layers.Input(shape=(80, 80, 4,))

    # First Conv2D layer
```

```

layer1 = layers.Conv2D(filters=32, kernel_size=8, strides=4, activation='relu',
                        kernel_initializer=initializers.TruncatedNormal(),
                        bias_initializer=initializers.Constant(0.01))

# MaxPooling layer
layer2 = layers.MaxPool2D(pool_size=2, strides=2, padding="same")(layer1)

# Second Conv2D layer
layer3 = layers.Conv2D(filters=64, kernel_size=3, strides=2, activation='relu',
                        kernel_initializer=initializers.TruncatedNormal(),
                        bias_initializer=initializers.Constant(0.01))

# Flatten layer
layer4 = layers.Flatten()(layer3)

# Dense layer for output
q_value = layers.Dense(units=2, activation="linear",
                        kernel_initializer=initializers.TruncatedNormal(),
                        bias_initializer=initializers.Constant(0.01))

return keras.Model(inputs=state, outputs=q_value)

```

Plot the model summary to make sure that the network is the same as expected.

```
In [43]: model = create_q_model()
print(model.summary())
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_9 (InputLayer)	[(None, 80, 80, 4)]	0
conv2d_8 (Conv2D)	(None, 19, 19, 32)	8224
max_pooling2d (MaxPooling2D)	(None, 10, 10, 32)	0
conv2d_9 (Conv2D)	(None, 4, 4, 64)	18496
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 2)	2050
=====		
Total params: 28,770		
Trainable params: 28,770		
Non-trainable params: 0		

None

```

2024-11-15 00:05:55.212783: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2 AVX AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
2024-11-15 00:05:55.213400: I tensorflow/core/common_runtime/process_util.cc:146] Creating new thread pool with default inter op setting: 2. Tune using inter_op_parallelism_threads for best performance.

```

## Deep Q-learning

We're now ready to implement the Q-learning algorithm. There are some subtle details in the implementation that you need to sort out. First, recall that the update rule for Q learning is

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma \cdot \max_{a'} Q(\text{next}(s, a), a') - Q(s, a))$$

where  $\gamma$  is the discount factor and  $\alpha$  can be viewed as the step size or learning rate for gradient ascent.

We'll set these as follows:

```
In [44]: gamma = 0.99          # decay rate of past observations
         step_size = 1e-4      # step size
```

## Estimation with experience replay

At the beginning of training, we spend 10,000 steps taking random actions, as a means of observing the environment.

We build a replay memory of length 10,000 steps, and every time we update the weights of the network, we sample a batch of size 32 and perform a Q-learning update on this batch.

After we have collected 10,000 steps of new data, we discard the old data, and replace it with the new "experiences."

```
In [45]: observe = 10000      # timesteps to observe before training
         replay_memory = 10000 # number of previous transitions to remember
         batch_size = 32       # size of each batch
```

## 3.3 Justify the data collection

Why does it make sense to maintain the replay memory of a fixed size instead of including all of the historical data?

maintaining a fixed-size replay memory is a practical choice to keep memory requirements manageable, reduce temporal correlation in training data, adapt the learning to the latest policy, and prevent the agent from being influenced by outdated and less relevant data.

1. **Memory Constraints:** Storing all historical data can become computationally infeasible for practical problems, especially as training goes on for many steps. Deep reinforcement learning can require millions of time steps, and maintaining all that data would require an enormous amount of memory. Fixing the size of the replay memory helps keep memory usage manageable.
2. **Efficiency in Learning:** As the model learns, older data becomes less useful or

even counterproductive, especially in a non-stationary environment where the dynamics or optimal policy may change over time. Including all data would increase the training time while introducing potentially outdated data that could slow down learning or lead to instability.

3. **Avoiding Correlation:** Experience replay aims to break the temporal correlation between consecutive experiences by shuffling and randomly sampling from the replay memory. This helps reduce the chance that the neural network overfits to sequences in the data. A fixed replay memory size ensures that the data in memory is sufficiently diverse and does not include long, correlated sequences from the past.
4. **Better Adaptation to Current Dynamics:** Since the agent's policy is constantly evolving during training, the relevance of older experiences decreases over time. Using a fixed-size buffer with the most recent experiences allows the agent to adapt more effectively to its current policy and environment dynamics, rather than being influenced by potentially outdated and less relevant experiences.

## Exploration vs exploitation

When performing Q-learning, we face the tradeoff between exploration and exploitation. To encourage exploration, a simple strategy is to take a random action at each step with certain probability.

More precisely, for each time step  $t$  and state  $s_t$ , with probability  $\epsilon$ , the algorithm takes a random action (wing flap or do nothing), and with probability  $1 - \epsilon$  the algorithm takes a greedy action according to  $a_t = \arg \max_a Q_\theta(s_t, a)$ . Here  $\theta$  refers to the parameters of our CNN.

```
In [46]: # value of epsilon
epsilon = 0.05
```

## 3.4 Complete the Q-learning algorithm

Next you will need to complete the Q-learning algorithm by filling in the missing code in the following function. The missing parts include

- Taking a greedy action
- Given a batch of samples  $\{(s_t, a_t, r_t, s_{t+1}, \text{terminal}_t)\}_{t \in B}$ , computing the corresponding  $Q_\theta(s_t, a_t)$ .
- Given a batch of samples  $\{(s_t, a_t, r_t, s_{t+1}, \text{terminal}_t)\}_{t \in B}$ , computing the corresponding updated Q-values

$$\hat{y}(s_t, a_t) = \begin{cases} r_t + \gamma \max_a Q_\theta(s_{t+1}, a), & \text{if } \text{terminal}_t = 0, \\ r_t, & \text{otherwise.} \end{cases}$$

Then, the mean squared error loss for the batch is



$$\frac{1}{|B|} \sum_{t \in B} (\hat{y}(s_t, a_t) - Q_{\theta}(s_t, a_t))^2.$$

You're now ready to play the game! Just run the cell below; do not change the code.

```
In [53]: def dql_flappy_bird(model, optimizer, loss_function):
    # initiate a game
    game = flappy_bird.GameState()

    # store the previous state, action and transitions
    history_data = deque()

    # get the first observation by doing nothing and preprocess the image
    do_nothing = np.zeros(num_actions)
    do_nothing[0] = 1
    image, reward, terminal = game.frame_step(do_nothing)

    # preprocess to get the state
    current_state = preprocess(image_raw=image)

    # training
    t = 0

    while t < 50000:
        if epsilon > np.random.rand(1)[0]:
            # random action
            action = np.random.choice(num_actions)
        else:
            # compute the Q function
            current_state_tensor = tf.convert_to_tensor(current_state)
            current_state_tensor = tf.expand_dims(current_state_tensor, 0)
            q_value = model(current_state_tensor, training=False)

            # greedy action
            action = np.argmax(q_value.numpy())

        # take the action and observe the reward and the next state
        action_vec = np.zeros([num_actions])
        action_vec[action] = 1
        image_raw, reward, terminal = game.frame_step(action_vec)
        next_state = preprocess(current_state=current_state,
                                image_raw=image_raw)

        # store the observation
        history_data.append((current_state, action, reward, next_state,
                             terminal))
        if len(history_data) > replay_memory:
            history_data.popleft() # discard old data

        # train if done observing
        if t > observe:

            # sample a batch
            batch = random.sample(history_data, batch_size)
            state_sample = np.array([d[0] for d in batch])
            action_sample = np.array([d[1] for d in batch])
            reward_sample = np.array([d[2] for d in batch])
```

```

state_next_sample = np.array([d[3] for d in batch])
terminal_sample = np.array([d[4] for d in batch])

# compute the updated Q-values for the samples
next_q_values = model(tf.convert_to_tensor(state_next_sample))
max_next_q_values = np.max(next_q_values, axis=1)
updated_q_values = reward_sample + (1 - terminal_sample) * gamma

# train the model on the states and updated Q-values
with tf.GradientTape() as tape:
    # compute the current Q-values for the samples
    current_q_values = model(tf.convert_to_tensor(state_sample))
    current_q_values = tf.reduce_sum(current_q_values * tf.one_hot(
        actions, num_actions))

    # compute the loss
    loss = loss_function(updated_q_values, current_q_values)

    # backpropagation
    grads = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(grads, model.trainable_variables))
else:
    loss = 0

# update current state and counter
current_state = next_state
t += 1

# print info every 500 steps
if t % 500 == 0:
    print(f"STEP {t} | PHASE {'observe' if t <= observe else 'train'} | ACTION {action} | REWARD {reward} | LOSS {loss}")

```

```

In [55]: def playgame(start_from_ckpt=False, ckpt_path=None):

    #! DO NOT change the random seed !
    np.random.seed(4)

    if start_from_ckpt:
        # if you want to start from a checkpoint
        model = keras.models.load_model('ckpt_path')
    else:
        model = create_q_model()

    # specify the optimizer and loss function
    optimizer = keras.optimizers.Adam(learning_rate=step_size, clipnorm=1)
    loss_function = keras.losses.MeanSquaredError()

    # play the game
    dql_flappy_bird(model=model, optimizer=optimizer, loss_function=loss_

```

```

In [56]: playgame()

```

STEP 500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 1000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 1500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 2000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 2500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 3000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 3500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 4000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 4500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 5000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 5500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 6000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 6500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 7000		PHASE observe		ACTION 0		REWARD 0.1		LOSS 0
STEP 7500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 8000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 8500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 9000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 9500		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 10000		PHASE observe		ACTION 1		REWARD 0.1		LOSS 0
STEP 10500		PHASE train		ACTION 1		REWARD -1		LOSS 0.01683366112411022
2								
STEP 11000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0261604916304349
9								
STEP 11500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0032830291893333
197								
STEP 12000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0105573590844869
61								
STEP 12500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.1992755681276321
4								
STEP 13000		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.3192921280860901
STEP 13500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0055427094921469
69								
STEP 14000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0120961070060729
98								
STEP 14500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0934694781899452
2								
STEP 15000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0077452268451452
255								
STEP 15500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0255564786493778
23								
STEP 16000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0023611634969711
304								
STEP 16500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0032806308008730
41								
STEP 17000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0156063009053468
7								
STEP 17500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0031164716929197
31								
STEP 18000		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0058163255453109
74								
STEP 18500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0053410101681947
71								
STEP 19000		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0084379604086279
87								
STEP 19500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0225962251424789
43								
STEP 20000		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0354456529021263
1								
STEP 20500		PHASE train		ACTION 0		REWARD 1		LOSS 0.07592210173606873

STEP 21000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.1060530617833137
5								
STEP 21500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0451483316719532
STEP 22000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0233204104006290
44								
STEP 22500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0368303284049034
1								
STEP 23000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0492456592619419
1								
STEP 23500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.2079674005508422
9								
STEP 24000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0197287499904632
57								
STEP 24500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0295626968145370
5								
STEP 25000		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0230500735342502
6								
STEP 25500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0272255428135395
05								
STEP 26000		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0161113664507865
9								
STEP 26500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0409014299511909
5								
STEP 27000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0776953771710395
8								
STEP 27500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0688210725784301
8								
STEP 28000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0387376807630062
1								
STEP 28500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0679298043251037
6								
STEP 29000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.1000777035951614
4								
STEP 29500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0998431742191314
7								
STEP 30000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0406733676791191
1								
STEP 30500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.5376922488212585
STEP 31000		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.0564160719513893
1								
STEP 31500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0664462447166442
9								
STEP 32000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0997502505779266
4								
STEP 32500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.4176612496376037
6								
STEP 33000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.1167316064238548
3								
STEP 33500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.3993814885616302
5								
STEP 34000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.1572796106338501
STEP 34500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.2133585810661316
STEP 35000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.2799621820449829
STEP 35500		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.0748222470283508
3								
STEP 36000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.2792862951755523
7								
STEP 36500		PHASE train		ACTION 1		REWARD 0.1		LOSS 0.1790128946304321
3								
STEP 37000		PHASE train		ACTION 0		REWARD 0.1		LOSS 0.3140020370483398

```

4
STEP 37500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.1833860278129577
6
STEP 38000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 1.579838752746582
STEP 38500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.0837092474102974
STEP 39000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.2358281016349792
5
STEP 39500 | PHASE train | ACTION 1 | REWARD 0.1 | LOSS 0.1793413311243057
3
STEP 40000 | PHASE train | ACTION 1 | REWARD 0.1 | LOSS 0.6232079267501831
STEP 40500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.2786513268947601
3
STEP 41000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.1860232502222061
2
STEP 41500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.2037529796361923
2
STEP 42000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.4208776950836181
6
STEP 42500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.7783730030059814
STEP 43000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 1.1934220790863037
STEP 43500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.2332394123077392
6
STEP 44000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.3275186419486999
5
STEP 44500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.3040559887886047
4
STEP 45000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.2225796580314636
2
STEP 45500 | PHASE train | ACTION 1 | REWARD 0.1 | LOSS 0.3317279517650604
STEP 46000 | PHASE train | ACTION 1 | REWARD 0.1 | LOSS 0.1620420515537262
STEP 46500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.8455294370651245
STEP 47000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.1906527727842331
STEP 47500 | PHASE train | ACTION 1 | REWARD 0.1 | LOSS 0.2109429985284805
3
STEP 48000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.3969378173351288
STEP 48500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.2077260315418243
4
STEP 49000 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.2027528136968612
7
STEP 49500 | PHASE train | ACTION 0 | REWARD 0.1 | LOSS 0.152142733335495
STEP 50000 | PHASE train | ACTION 1 | REWARD 0.1 | LOSS 0.3111290335655212
4

```

### 3.5 Describe the training

Describe what you see by answering the following questions:

- In the early stage of training (within 2,000 steps in the *explore* phase), describe the behavior of the Flappy Bird. What do you think is the greedy policy given by the estimation of the Q-function in this stage?
- Describe what you see after roughly 5,000 training steps. Do you see any improvement? In particular, compare Flappy's behavior with their behavior in the early stages of training.
- Explain why the performance has improved, by relating to the model design such as the replay memory and the exploration.

# Analysis of the Q-learning Training Process for Flappy Bird

## 1. Early Stage of Training (Within 2,000 Steps in the Observe Phase)

- **Behavior of Flappy Bird:** In the early phase (first 10,000 steps, specifically within 2,000 steps), the agent is primarily in the observation phase, where it takes random actions to gather diverse data about the environment. During this stage, the agent's actions are not guided by any understanding of the game or optimal strategies.
  - **Greedy Policy Estimate:** Since the Q-function is still largely untrained, the model's value estimates are quite arbitrary. The Q-values are essentially uninformative at this point, which means the agent is unlikely to have developed a coherent greedy policy. The actions taken are random, resulting in exploration without a clear strategy to avoid obstacles or navigate effectively. This is why the behavior appears inconsistent and non-optimal.

## 2. Behavior after Roughly 5,000 Training Steps

- **Improvement in Performance:** After 5,000 training steps, the model has started transitioning into the training phase and is updating its weights based on experiences sampled from the replay memory. The loss starts to decrease, indicating that the model is learning to approximate the optimal Q-values better.
  - **Comparison to Early Stages:** Compared to the early stages, where the agent was exploring randomly, there is now some evidence that the agent is learning to associate certain states with higher expected rewards. This results in more structured and purposeful actions that improve the game performance (e.g., avoiding obstacles more effectively or maintaining a better altitude).

## 3. Reasons for Improved Performance

- **Replay Memory:** The replay memory helps reduce the correlation between consecutive samples by shuffling experiences and using them for training in batches. This stabilizes learning because the agent is not overfitting to a specific sequence of events, and the model can learn from a diverse range of experiences. The fixed replay memory ensures that only relatively recent experiences are used, making the training data more relevant to the agent's current policy.
- **Exploration vs. Exploitation:** In the beginning, the agent explores the environment using random actions (`epsilon` is high). This allows the agent to gather varied experiences that cover different states of the environment. Over time, as `epsilon` decreases, the agent shifts from exploration to exploitation. This allows the agent to act according to the learned Q-values more frequently, thereby following a policy that is becoming increasingly optimal as training progresses. This balance between exploration and exploitation is critical for discovering an effective policy.

- **Batch Updates for Stability:** Using batch updates also contributes to stable learning. Instead of updating the model after every step (which could lead to unstable learning due to high variance in rewards), the agent samples a batch of experiences, effectively averaging the updates over a more comprehensive set of experiences. This helps in reducing variance and improving the stability of Q-value updates.
- **Q-learning Update Rule:** The Q-learning update rule incorporates the immediate reward and the maximum estimated future reward, weighted by the discount factor (  $\gamma$  ). Initially, the Q-values are arbitrary, but as training progresses, these values become more accurate estimates of the expected cumulative reward for each state-action pair. The combination of current rewards and future expected rewards helps the agent learn sequences of actions that yield high long-term gains rather than being myopically focused on short-term rewards.

Overall, the combination of replay memory, exploration-exploitation trade-off, and the Q-learning update mechanism allows the agent to learn effective policies over time, improving its performance in the game. This is evident from the observed reduction in loss and improvement in game behavior as the number of training steps increases.

It takes a long time to fully train the network, so you're not required to complete the training. Here's a [video](#) showing the performance of a well trained DQN.